

Week 3: Stream (1)


Instructor: Daejin Choi (djchoi@inu.ac.kr)



INCHEON
NATIONAL
UNIVERSITY

Static vs. Stream Data

- REMIND: Big data means not only “size”, but “velocity”



Characteristics of **Stream** Data

- Infinite, Burst, Non-stationary
 - **We do not know the entire data set in advance**
- only **INSTANTLY** accessible
- **Stream Management** is important when the input rate is controlled **externally**:
 - Google queries
 - Twitter or Facebook status updates

Applications

- **Mining query streams**

- Google wants to know what queries are more frequent today than yesterday

- **Mining click streams**

- Yahoo wants to know which of its pages are getting an unusual number of hits in the past hour

- **Mining social network news feeds**

- E.g., look for trending topics on Twitter, Facebook

- **Sensor Networks**

- Many sensors feeding into a central controller

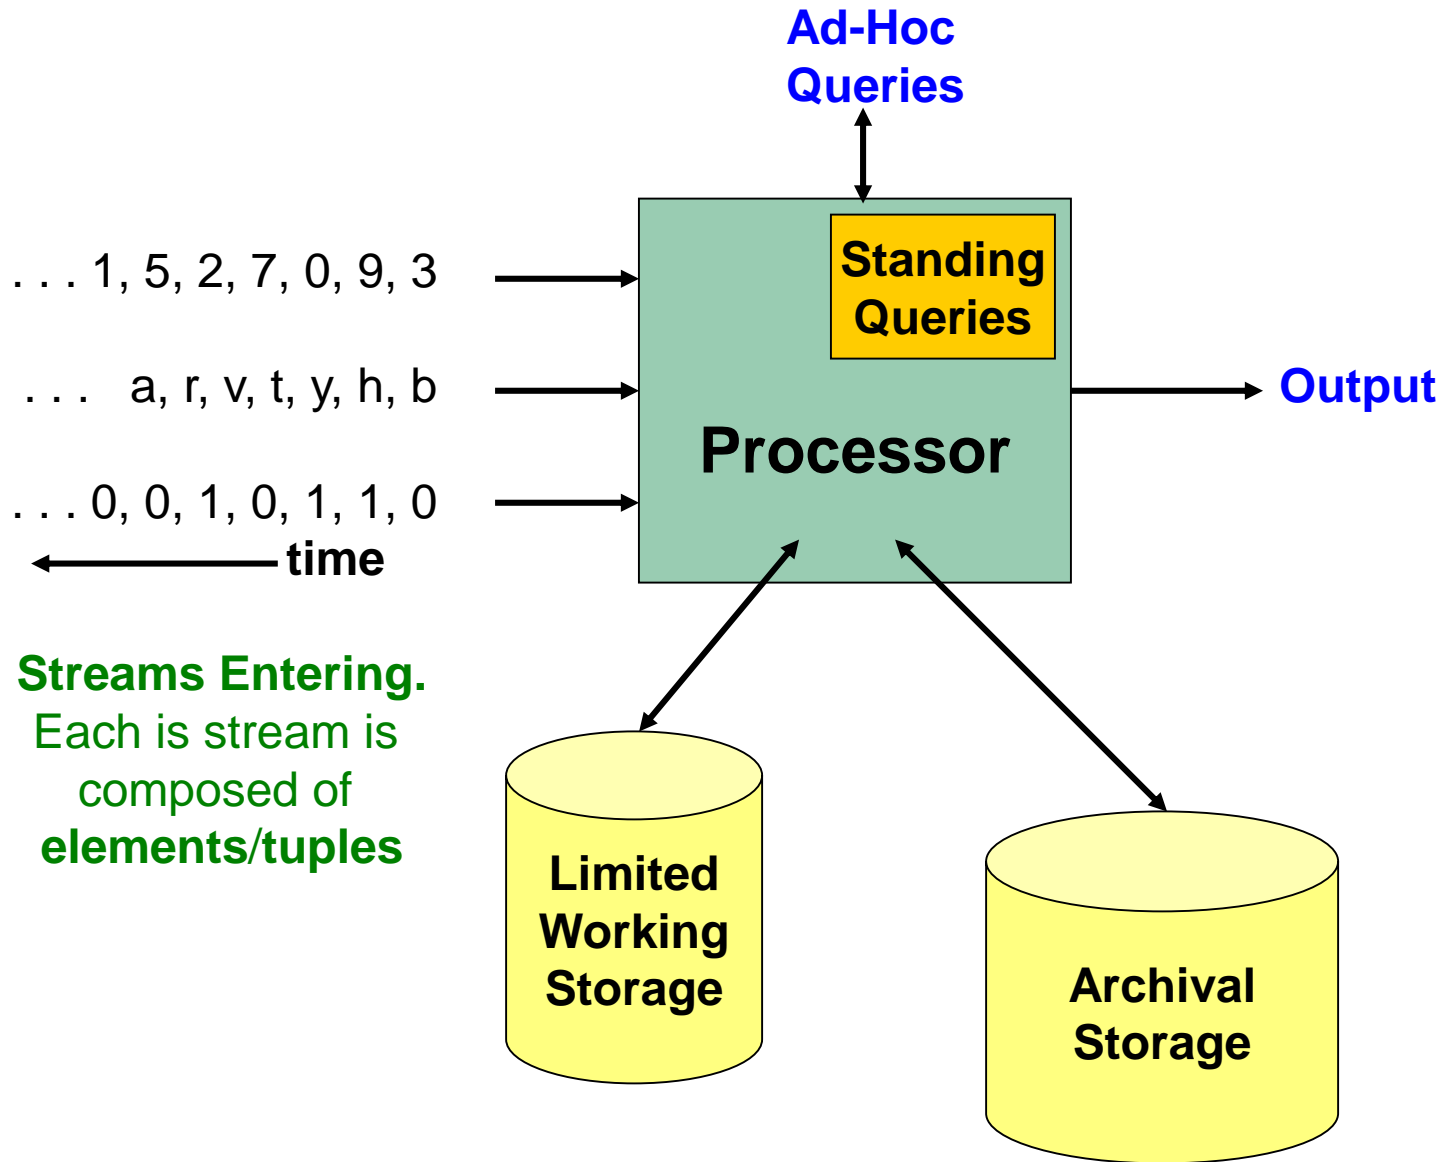
- **IP packets monitored at a switch**

- Gather information for optimal routing
- Detect denial-of-service attacks

The Stream Model

- Input **elements** enter at a *rapid* rate, at one or more input ports (i.e., **streams**)
 - **We call elements of the stream tuples**
- The system **cannot** store the entire stream
- **Q:** How do you make critical calculations about the stream using a **limited amount of (secondary) memory?**

General Stream Processing Model



Side note: Streaming Alg. in Machine Learning?

- **Online Learning** enables a machine learning model to continuously learn from the **recent data stream**
 - An algorithm to learn from it and slowly adapt to the changes in data
- **Example: Stochastic Gradient Descent (SGD)**
- **Idea: Do slow updates to the model**
 - **SGD** (SVM, Perceptron) makes small updates
 - **So:** First train the classifier on training data.
 - **Then:** For every example from the stream, we slightly update the model (using small learning rate)

Operations on Data Streams

- In conclusion, we have to choose **a subset of input streams**

- **Sampling data from a stream**

- Construct a random sample

- **Queries over sliding windows**

- Number of items of type x in the last k elements of the stream

- **Filtering a data stream**

- Select elements with property x from the stream

- **Counting distinct elements**

- Number of distinct elements in the last k elements of the stream

- **Estimating moments**

- Estimate avg./std. dev. of last k elements

- ...

Week 3

Week 4

Let's start from sampling!

- Sample a **fixed proportion** of elements in the stream (say 1 in 10)
- Maintain a **random sample of fixed size** over a potentially infinite stream
 - At any “time” k we would like a random sample of s elements
 - **What is the property of the sample we want to maintain?**
 - : For all time steps k , each element seen so far has equal prob. of being sampled

Sampling from a Data Stream: **Sampling a fixed proportion**

Sampling a Fixed Proportion

- **Scenario:** Search engine query stream
 - **Stream of tuples:** (user, query, time)
 - **Answer questions such as: How often did a user run the same query in a single day**
 - Have space to store **$1/10^{\text{th}}$** of query stream
- **Naïve solution:**
 - Generate a random integer in **$[0..9]$** for each query
 - Store the query if the integer is **0**, otherwise discard

Problem with Naïve Approach

- **Simple question: What fraction of queries by an user are duplicates?**
 - Suppose each user issues x queries once and d queries twice (total of $x+2d$ queries)
 - **Correct answer:** $d/(x+d)$
 - **Proposed solution: We keep 10% of the queries**
 - Sample will contain $x/10$ of the singleton queries and $2d/10$ of the duplicate queries at least once
 - But only $d/100$ pairs of duplicates
 - $d/100 = 1/10 \cdot 1/10 \cdot d$
 - Of d "duplicates" $18d/100$ appear exactly once
 - $18d/100 = ((1/10 \cdot 9/10) + (9/10 \cdot 1/10)) \cdot d$
 - **So the sample-based answer is** $\frac{\frac{x}{10} + \frac{d}{100} + \frac{18d}{100}}{\frac{x}{10} + \frac{d}{100} + \frac{18d}{100}} = \frac{d}{10x+19d}$

How to solve?

- Where the error comes from? → “Probability of selection”
- A possible solution? → Sampling user, not query
 - Pick **1/10th** of **users** and take all their searches in the sample
 - Use a hash function that hashes the user name or user id uniformly into 10 buckets

Generalized Solution

- **Stream of tuples with keys:**
 - Key is a subset of each tuple's components
 - e.g., tuple is (user, search, time); key is **user**
 - Choice of key depends on application
- **To get a sample of a/b fraction of the stream:**
 - Hash each tuple's key uniformly into **b** buckets
 - Pick the tuple if its hash value is at most **a**



Hash table with **b** buckets, pick the tuple if its **hash** value is at most **a** .

How to generate a 30% sample?

Hash into $b=10$ buckets, take the tuple if it hashes to one of the first 3 buckets
(Assumption: Hash distributes users “uniformly”)

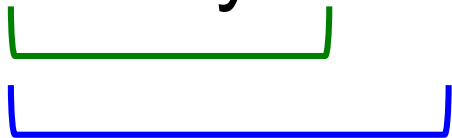
Sampling from a Data Stream:

Sampling a fixed-size tuples

Maintaining a fixed-size sample

- **Suppose we need to maintain a random sample S of size exactly s tuples**
 - E.g., main memory size constraint
- **Why?** Don't know length of stream in advance
- **Suppose at time n we have seen s items**
 - **Goal: Each item is in the sample S with equal prob. s/n**

Illustration ($s=2$)

- **Stream:** a x c y z k c d e g...

- At $n=5$, each of the first 5 tuples is included in the sample \mathbf{S} with equal prob.
- At $n=7$, each of the first 7 tuples is included in the sample \mathbf{S} with equal prob.
- **Impractical solution would be to store all the n tuples seen so far and out of them pick s at random**

Solution: Reservoir Sampling

■ Algorithm

- Store all the first s elements of the stream to S
- Suppose we have seen $n-1$ elements, and now the n^{th} element arrives ($n > s$)
 - With probability s/n , keep the n^{th} element, else discard it
 - If we picked the n^{th} element, then it replaces one of the s elements in the sample S , picked uniformly at random

■ Claim: This algorithm maintains a sample S with the desired property:

- After n elements, the sample contains each element seen so far with probability s/n

Proof of Reservoir Sampling By Induction

- Assume that after n elements, the sample contains each element seen so far with probability s/n
- We need to show that after seeing element $n+1$ the sample maintains the property
 - Sample contains each element seen so far with probability $s/(n+1)$
- **Base case:**
 - After we see $n=s$ elements the sample S has the desired property
 - Each out of $n=s$ elements is in the sample with probability $s/s = 1$

Proof (cont'd)

- **Inductive hypothesis:** After n elements, the sample S contains each element seen so far with prob. s/n
- **Now element $n+1$ arrives**
- **Inductive step:** For elements already in S , probability that the algorithm keeps it in S is:

$$\underbrace{\left(1 - \frac{s}{n+1}\right)}_{\text{Element } n+1 \text{ discarded}} + \underbrace{\left(\frac{s}{n+1}\right)}_{\text{Element } n+1 \text{ not discarded}} \underbrace{\left(\frac{s-1}{s}\right)}_{\text{Element in the sample not picked}} = \frac{n}{n+1}$$

- So, at time n , tuples in S were there with prob. s/n
- Time $n \rightarrow n+1$, tuple stayed in S with prob. $n/(n+1)$
- So prob. tuple is in S at time $n+1 = \frac{s}{n} \cdot \frac{n}{n+1} = \frac{s}{n+1}$

Queries over a (long) Sliding Window

Sliding Windows

- A useful model of stream processing is that queries are about a *window* of length N , the N most recent elements received
- **Interesting case:** N is so large that the data cannot be stored in memory, or even on disk
 - Or, there are so many streams that windows for all cannot be stored
 - The estimation is required. We will finally cover later!

Sliding Window: 1 Stream

- **Sliding window on a single stream:**

N = 6

q w e r t y u i o p a s d f g h j k l z x c v b n m

q w e r t y u i o p a s d f g h j k l z x c v b n m

q w e r t y u i o p a s d f g h j k l z x c v b n m

q w e r t y u i o p a s d f g h j k l z x c v b n m

← Past

Future →

Example Problem: Counting Bits (1)


■ Problem:

- Given a stream of **0**s and **1**s
- Be prepared to answer queries of the form
How many 1s are in the last k bits? where $k \leq N$

■ Obvious solution:

Store the most recent N bits

- When new bit comes in, discard the $N+1^{\text{st}}$ bit



Suppose $N=6$


Example Problem: Counting Bits (2)

- You **CANNOT** get an exact answer without storing the entire window

- **Real Problem:**

What if we cannot afford to store N bits?


- **E.g.**, we're processing 1 billion streams and **$N = 1$ billion**



- But we are happy with **an approximate answer** inferred from **the summarized results**

An attempt: Simple solution

- **Q: How many 1s are in the last N bits?**
- A simple solution that does not really solve our problem: **Uniformity assumption**

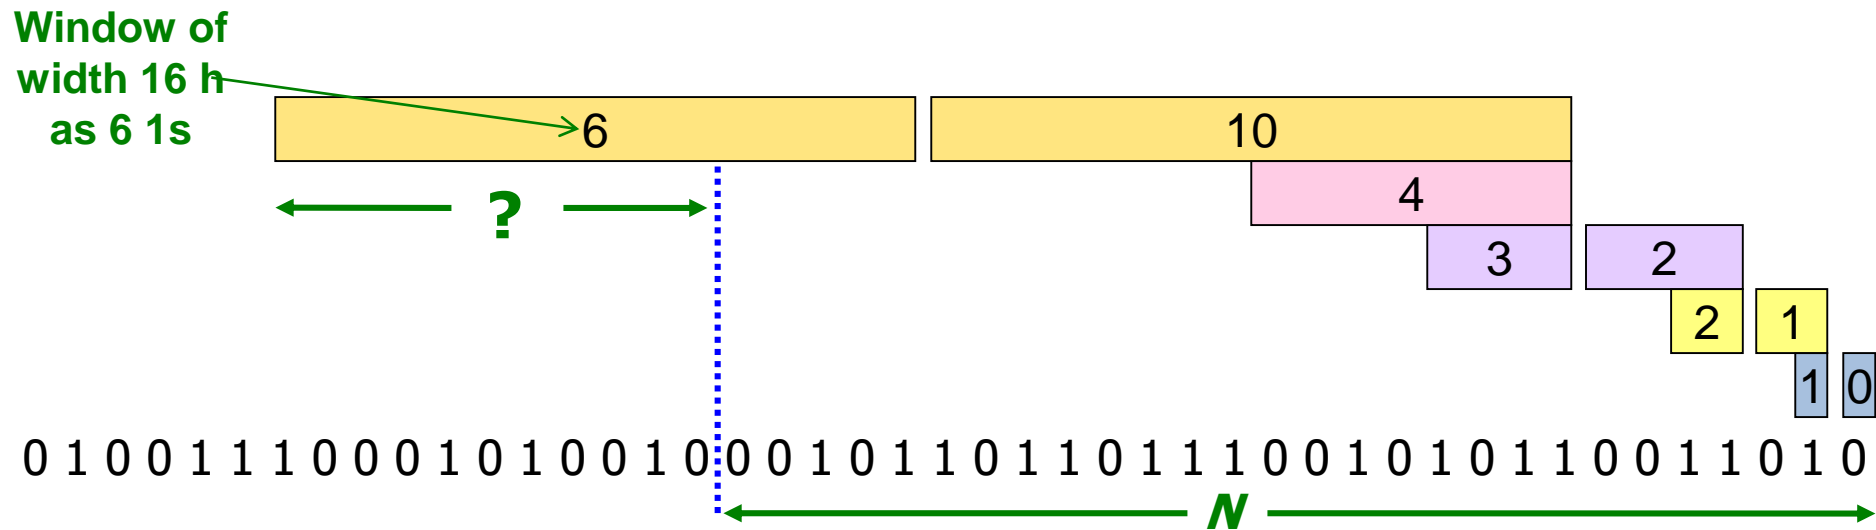


- **Maintain 2 counters:**
 - **S** : number of 1s from the beginning of the stream
 - **Z** : number of 0s from the beginning of the stream
- **How many 1s are in the last N bits?** $N \cdot \frac{S}{S+Z}$
- **But, what if stream is non-uniform?**
 - What if distribution changes over time?

- **DGIM solution that does not assume uniformity**
- We store $O(\log^2 N)$ bits per stream
- **Solution gives approximate answer, never off by more than 50%**
 - Error factor can be reduced to any fraction > 0 , with more complicated algorithm and proportionally more stored bits

Beginning Idea: Exponential Windows


- Summarize **exponentially increasing** regions of the stream, looking backward
- Drop small regions if they begin at the same point as a larger region




We can reconstruct the count of the last N bits, except we are not sure how many of the last 6 1s are included in the N

What's Good?

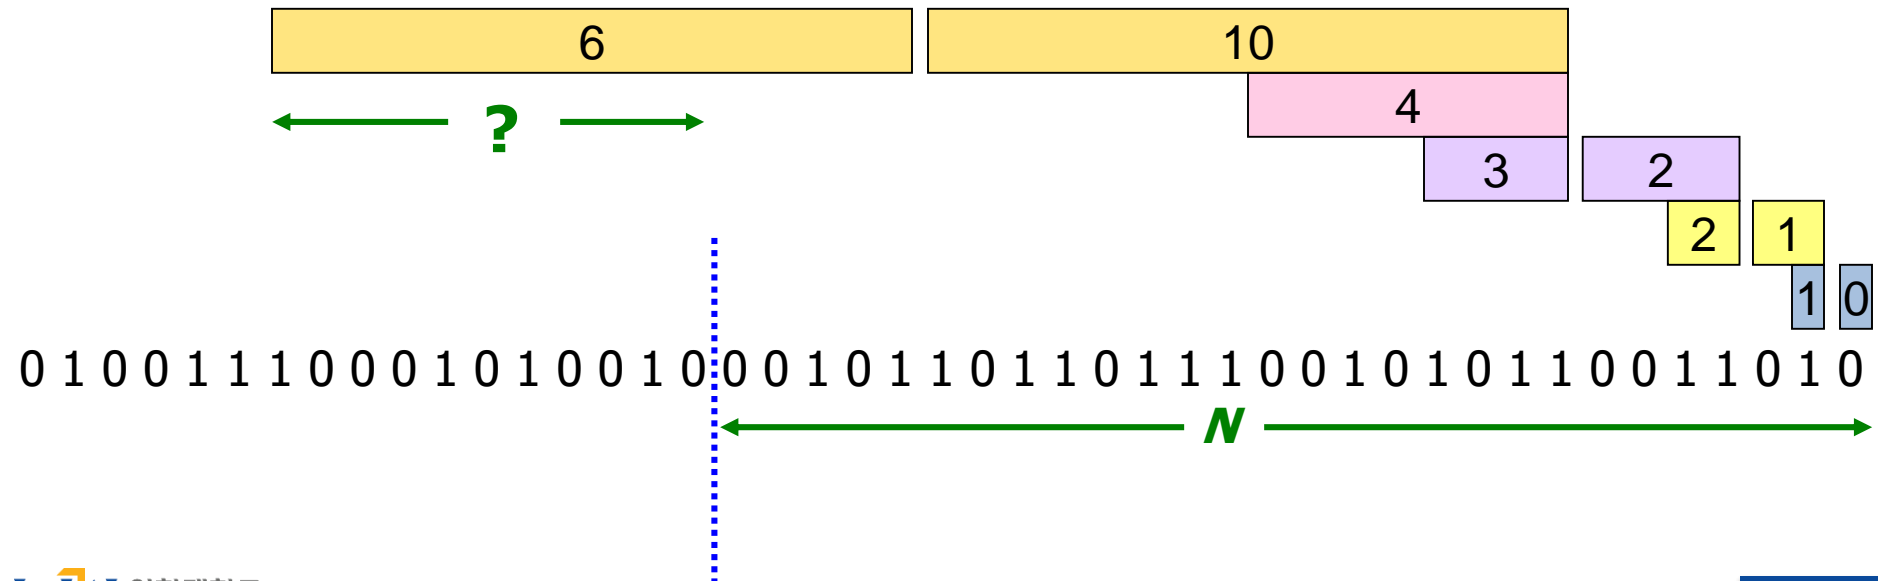
- **Stores only $O(\log^2 N)$ bits**
 - $O(\log N)$ counts of $\log_2 N$ bits each


Window counts



Window size
- **Easy update as more bits enter**
- Error in count no greater than the number of **1s** in the “**unknown**” area

What's Not So Good?

- As long as the **1s** are fairly evenly distributed, the error due to the unknown region is small – **no more than 50%**
- But it could be that all the **1s** are in the unknown area at the end
- In that case, **the error changes w.r.t. N !**



- **Intuition:** We do not have to count both 0s and 1s
- **Idea:** Instead of summarizing fixed-length blocks, summarize **blocks with specific number of 1s**:
 - Let the block *sizes* (number of **1s**) increase exponentially




- When there are few 1s in the window, block sizes stay small, so errors are small

DGIM: Timestamps

- Each bit in the stream has a *timestamp*, starting **1, 2, ...**
- Record timestamps modulo **N** (**the window size**), so we can represent any **relevant** timestamp in $O(\log_2 N)$ bits
 - Store **$N = 4$** using 2 bits

DGIM: Buckets


- A **bucket** in the DGIM method is a record consisting of:
 - **(A)** The timestamp of its end [$O(\log N)$ bits]
 - **(B)** The number of 1s between its beginning and end [$O(\log \log N)$ bits]
- **Constraint on buckets:**
Number of **1s** must be a power of 2
 - That explains the $O(\log \log N)$ in **(B)** above



More Constraints

- Either **one** or **two** buckets with the same **power-of-2 number of 1s**
- **Buckets do not overlap in timestamps**
- **Buckets are sorted by size**
 - Earlier buckets are not smaller than later buckets
- Buckets disappear when their end-time is $> N$ time units in the past

Example: Representing a Stream by Buckets



Three properties of buckets that are maintained:

- Either **one** or **two** buckets with the same **power-of-2** number of **1s**
- Buckets do not overlap in timestamps
- Buckets are sorted by size

Updating Buckets (1)

- When a new bit comes in, drop the last (oldest) bucket if its end-time is prior to **N** time units before the current time
- **2 cases:** Current bit is **0** or **1**
- **If the current bit is 0:**
no other changes are needed

Updating Buckets (2)

- **If the current bit is 1:**

(1) Create a new bucket of size **1**, for just this bit

- **End timestamp = current time**

(2) If there are now **three buckets of size 1**, **combine the oldest two into a bucket of size 2**

(3) If there are now **three buckets of size 2**, **combine the oldest two into a bucket of size 4**

(4) **And so on ...**

Example: Updating Buckets

Current state of the stream:

10010101100010110101010101010110101010101011101010101110101011101010100010110010

Bit of value 1 arrives

00101011000101101010101010101101010101010111010101011101010111010101000101100101

Two orange buckets get merged into a yellow bucket

00101011000101101010101010101101010101010111010101011101010111010101000101100101

Next bit 1 arrives, new orange bucket is created, then 0 comes, then 1:

010110001011010101010101010110101010101110101010111010101000101100101101

Buckets get merged...

010110001011010101010101010110101010101110101010111010101000101100101101


State of the buckets after merging

010110001011010101010101010110101010101110101010111010101000101100101101

Return to Our Original Question


- **To estimate the number of 1s in the most recent N bits:**
 1. **Sum the sizes of all buckets but the last**
(note “size” means the number of 1s in the bucket)
 2. **Add half the size of the last bucket**
- **Remember:** We do not know how many **1s** of the last bucket are still within the wanted window

Example: Bucketized Stream



How can the error be bounded?

- Why is error 50%? Let's prove it!




- Suppose the last bucket has size 2^r
- Then by assuming 2^{r-1} (i.e., half) of its **1s** are still within the window, we make an error of at most 2^{r-1}
- Since there is at least one bucket of each of the sizes less than 2^r , the true sum is at least $1 + 2 + 4 + \dots + 2^{r-1} = 2^r - 1$
- Thus, error at most 50%

Further Reducing the Error

- Instead of maintaining **1** or **2** of each size bucket, we allow either **$r-1$** or **r** buckets (**$r > 2$**)
 - Except for the largest size buckets; we can have any number between **1** and **r** of those
- **Error is at most $O(1/r)$**
- By picking **r** appropriately, we can tradeoff between **number of bits we store** and **the error**

Extensions

- Can we use the same trick to answer queries **How many 1's in the last k ?** where $k < N$?
 - A:** Find earliest bucket **B** that overlaps with k .
Number of **1s** is the **sum of sizes of more recent buckets + 1/2 size of B**



Summary

- **Sampling a fixed proportion of a stream**
 - Sample size grows as the stream grows
- **Sampling a fixed-size sample**
 - Reservoir sampling
- **Counting the number of 1s in the last N elements**
 - Exponentially increasing windows → DGIM method

The background of the slide is an abstract geometric pattern composed of various shades of blue and light blue triangles and polygons, creating a low-poly or crystalline effect.

Thank you!

Instructor: Daejin Choi (djchoi@inu.ac.kr)